

Contents

| | |
|----------------------------------------------------------|-----------|
| 1 QSeaBattle “ Design Specification | 7 |
| 1.1 Purpose | 7 |
| 1.2 Scope and audience | 7 |
| 1.3 Structure of the specification | 8 |
| 1.4 Normative vs informative content | 8 |
| 1.5 Versioning and stability | 8 |
| 1.6 How to read this document | 9 |
| 1.7 Repository layout (informative) | 9 |
| 1.8 Design principle (non-negotiable) | 9 |
| 2 Algorithms overview (informative) | 9 |
| 2.1 Purpose | 9 |
| 2.2 1. Classical assisted algorithm (baseline) | 9 |
| 2.2.1 Problem | 9 |
| 2.2.2 High-level idea | 10 |
| 2.2.3 Algorithm (conceptual) | 10 |
| 2.3 2. Shared randomness algorithm | 10 |
| 2.3.1 Resource model | 10 |
| 2.3.2 Usage rules | 10 |
| 2.4 3. Deterministic baselines (Simple and Majority) . . | 10 |
| 2.4.1 Why they exist | 10 |
| 2.4.2 Simple baseline (conceptual) | 11 |
| 2.4.3 Majority baseline (conceptual) | 11 |
| 2.5 4. Lin trainable assisted algorithm | 11 |
| 2.5.1 Goal | 11 |
| 2.5.2 Algorithm | 11 |
| 2.6 5. Pyramid (Pyr) trainable assisted algorithm . . . | 11 |
| 2.6.1 Motivation | 11 |
| 2.6.2 Algorithm | 11 |
| 2.7 6. Training algorithms (summary) | 12 |
| 2.7.1 Imitation learning | 12 |
| 2.7.2 Key restriction | 12 |
| 2.8 Relationship to the rest of the spec | 12 |
| 2.9 Reader guidance | 12 |
| 3 Conventions | 12 |
| 3.1 Purpose | 12 |
| 3.2 Notation | 13 |
| 3.2.1 Sizes and shapes | 13 |
| 3.2.2 Binary values | 13 |
| 3.3 Terminology | 13 |
| 3.3.1 Game vs tournament | 13 |
| 3.3.2 Players and models | 13 |
| 3.3.3 Shared randomness (SR) | 13 |

| | |
|------------------------------------------------------------------------|-----------|
| 3.4 Determinism, seeding, and reproducibility | 14 |
| 3.5 Channel noise | 14 |
| 3.6 Contract language (RFC-style) | 14 |
| 3.7 Interface conventions | 14 |
| 3.7.1 decide(...) | 14 |
| 3.7.2 Dtypes | 14 |
| 3.8 Error handling | 14 |
| 3.9 Document status | 15 |
| 4 Class GameLayout | 15 |
| 4.1 Purpose | 15 |
| 4.2 Location | 15 |
| 4.3 Public attributes | 15 |
| 4.4 Public methods | 16 |
| 4.4.1 from_dict(parameters: dict) -> GameLayout . . | 16 |
| 4.4.2 to_dict() -> dict | 16 |
| 4.5 Invariants | 16 |
| 4.6 Failure modes | 16 |
| 5 Class GameEnv | 16 |
| 5.1 Purpose | 16 |
| 5.2 Location | 16 |
| 5.3 Public methods | 16 |
| 5.3.1 reset() -> None | 16 |
| 5.3.2 provide() -> tuple[np.ndarray, np.ndarray] . . | 16 |
| 5.3.3 evaluate(shoot: int) -> float | 16 |
| 5.3.4 apply_channel_noise(comm: np.ndarray) -> np.ndarray | 17 |
| 5.4 Invariants | 17 |
| 6 Class Game | 17 |
| 6.1 Purpose | 17 |
| 6.2 Method | 17 |
| 6.2.1 play() | 17 |
| 7 Class Tournament | 17 |
| 7.1 Purpose | 17 |
| 7.1.1 tournament() -> TournamentLog | 17 |
| 8 Class TournamentLog | 17 |
| 8.1 Purpose | 17 |
| 8.1.1 Methods | 17 |
| 9 Base player interfaces | 18 |
| 9.1 Purpose | 18 |
| 9.2 Class Players | 18 |
| 9.2.1 Purpose | 18 |

| | |
|---------------------------------------------------------|-----------|
| 9.2.2 Public methods | 18 |
| 9.2.3 Invariants | 18 |
| 9.3 Class PlayerA | 18 |
| 9.3.1 Purpose | 18 |
| 9.3.2 Method | 18 |
| 9.3.3 Contract | 18 |
| 9.3.4 Default baseline behavior (random) | 19 |
| 9.4 Class PlayerB | 19 |
| 9.4.1 Purpose | 19 |
| 9.4.2 Method | 19 |
| 9.4.3 Contract | 19 |
| 9.4.4 Default baseline behavior (random) | 19 |
| 9.5 Related | 19 |
| 10 Deterministic and baseline players | 19 |
| 10.1 Purpose | 19 |
| 10.2 Class SimplePlayers(Players) | 20 |
| 10.2.1 Purpose | 20 |
| 10.2.2 Behavioral contract | 20 |
| 10.3 Class SimplePlayerA(PlayerA) | 20 |
| 10.3.1 Purpose | 20 |
| 10.3.2 Algorithm (normative) | 20 |
| 10.3.3 Invariants | 20 |
| 10.4 Class SimplePlayerB(PlayerB) | 20 |
| 10.4.1 Purpose | 20 |
| 10.4.2 Algorithm (normative) | 21 |
| 10.4.3 Notes | 21 |
| 10.4.4 Invariants | 21 |
| 10.5 Class MajorityPlayers(Players) | 21 |
| 10.5.1 Purpose | 21 |
| 10.5.2 Behavioral contract | 21 |
| 10.6 Class MajorityPlayerA(PlayerA) | 21 |
| 10.6.1 Purpose | 21 |
| 10.6.2 Algorithm (normative) | 21 |
| 10.6.3 Invariants | 22 |
| 10.7 Class MajorityPlayerB(PlayerB) | 22 |
| 10.7.1 Purpose | 22 |
| 10.7.2 Algorithm (normative) | 22 |
| 10.7.3 Invariants | 22 |
| 10.8 Related | 22 |
| 11 Assisted players (shared-randomness assisted) | 23 |
| 11.1 Purpose | 23 |
| 11.2 Shared randomness dependency | 23 |
| 11.3 Class SharedRandomness() | 23 |
| 11.3.1 Purpose | 23 |

| | | |
|-------------|-------------------------------------------------------------------|-----------|
| 11.4 | Class AssistedPlayers(Players) | 23 |
| 11.4.1 | Purpose | 23 |
| 11.4.2 | Location | 23 |
| 11.4.3 | Parameters | 23 |
| 11.4.4 | Behavioral contract | 24 |
| 11.5 | Class AssistedPlayerA(PlayerA) | 24 |
| 11.5.1 | Purpose | 24 |
| 11.5.2 | Behavioral contract (normative) | 24 |
| 11.5.3 | Algorithmic outline (informative) | 24 |
| 11.6 | Class AssistedPlayerB(PlayerB) | 24 |
| 11.6.1 | Purpose | 24 |
| 11.6.2 | Behavioral contract (normative) | 24 |
| 11.6.3 | Algorithmic outline (informative) | 25 |
| 11.7 | Invariants | 25 |
| 12 | Neural players (unassisted) | 25 |
| 12.1 | Purpose | 25 |
| 12.2 | Characteristics | 25 |
| 12.3 | Behavioral constraints | 25 |
| 12.4 | Notes | 25 |
| 13 | Shared randomness resource (non-classical) | 26 |
| 13.1 | Purpose | 26 |
| 13.2 | Location | 26 |
| 13.3 | Parameters | 26 |
| 13.4 | Public interface (conceptual) | 26 |
| 13.4.1 | reset() -> None | 26 |
| 13.4.2 | measurement_a(measurement: array_like) -> np.ndarray | 26 |
| 13.4.3 | measurement_b(measurement: array_like) -> np.ndarray | 26 |
| 13.5 | Behavioral contract (normative) | 27 |
| 13.5.1 | Single-use per party | 27 |
| 13.5.2 | Two-measurement lifecycle | 27 |
| 13.5.3 | Output format | 27 |
| 13.5.4 | Non-classical correlation requirement | 27 |
| 13.5.5 | No information leakage | 28 |
| 13.6 | Failure modes | 28 |
| 13.7 | Notes (informative) | 28 |
| 14 | SharedRandomnessLayer (differentiable implementation) | 28 |
| 14.1 | Purpose | 28 |
| 14.2 | Location | 28 |
| 14.3 | Inputs / Outputs | 28 |
| 14.4 | Behavioral contract | 29 |
| 14.5 | Recommended modes (informative) | 29 |

| | |
|-------------------------------------------------------|-----------|
| 14.6 Failure modes | 29 |
| 15 Trainable assisted models â€” Linear (Lin) | 29 |
| 15.1 Purpose | 29 |
| 15.2 Scope | 29 |
| 15.3 Class LinTrainableAssistedModelA | 30 |
| 15.4 Purpose | 30 |
| 15.5 Location | 30 |
| 15.6 Inputs | 30 |
| 15.7 Outputs | 30 |
| 15.8 Internal state | 30 |
| 15.9 Behavioral contract | 30 |
| 15.10 Invariants | 30 |
| 15.11 Class LinTrainableAssistedModelB | 31 |
| 15.12 Purpose | 31 |
| 15.13 Location | 31 |
| 15.14 Inputs | 31 |
| 15.15 Outputs | 31 |
| 15.16 Behavioral contract | 31 |
| 15.17 Invariants | 31 |
| 15.18 Notes / rationale | 31 |
| 16 Trainable assisted models â€” Pyramid (Pyr) | 32 |
| 16.1 Purpose | 32 |
| 16.2 Conceptual overview | 32 |
| 16.3 Class PyrTrainableAssistedModelA | 32 |
| 16.4 Purpose | 32 |
| 16.5 Location | 32 |
| 16.6 Inputs | 32 |
| 16.7 Outputs | 33 |
| 16.8 Internal state | 33 |
| 16.9 Iterative structure (normative) | 33 |
| 16.10 Behavioral contract | 33 |
| 16.11 Invariants | 33 |
| 16.12 Failure modes | 34 |
| 16.13 Class PyrTrainableAssistedModelB | 34 |
| 16.14 Purpose | 34 |
| 16.15 Location | 34 |
| 16.16 Inputs | 34 |
| 16.17 Outputs | 34 |
| 16.18 Internal state | 34 |
| 16.19 Reconstruction procedure (normative) | 34 |
| 16.20 Behavioral contract | 35 |
| 16.21 Invariants | 35 |
| 16.22 Notes / rationale | 35 |

| | |
|---------------------------------------------------------------------------|-----------|
| 17 Imitation learning for assisted models | 35 |
| 17.1Purpose | 35 |
| 17.2Data generation | 35 |
| 17.3Behavioral contract | 35 |
| 17.4Losses | 35 |
| 17.5Invariants | 36 |
| 18 DRU / DIAL training | 36 |
| 18.1Purpose | 36 |
| 18.2Algorithm (informative) | 36 |
| 18.3Constraints | 36 |
| 18.4Notes | 36 |
| 19 Reinforcement learning | 36 |
| 19.1Purpose | 36 |
| 19.2Supported approaches | 36 |
| 19.3Constraints | 37 |
| 19.4Notes | 37 |
| 20 Training considerations for Pyramid (Pyr) models | 37 |
| 20.1Purpose | 37 |
| 20.2Training regimes | 37 |
| 20.3Behavioral constraints | 37 |
| 20.4Diagnostics (recommended) | 37 |
| 20.5Invariants | 37 |
| 21 Utility module: Data generation | 38 |
| 21.1Purpose | 38 |
| 21.2Scope | 38 |
| 21.3Core utilities | 38 |
| 21.3.1generate_measurement_dataset_a(layout, num_samples, seed) | 38 |
| 21.3.2generate_measurement_dataset_b(layout, num_samples, seed) | 38 |
| 21.4Invariants | 38 |
| 22 Utility module: Imitation training | 38 |
| 22.1Purpose | 38 |
| 22.2Core utilities | 39 |
| 22.2.1train_layer(layer, dataset, loss, epochs, metrics=None) | 39 |
| 22.2.2train_model(model, datasets) | 39 |
| 22.3Invariants | 39 |
| 23 Utility module: Layer and weight transfer | 39 |
| 23.1Purpose | 39 |
| 23.2Core utilities | 39 |

| | |
|----------------------------------------------------|-----------|
| 23.2.1transfer_assisted_model_a_layer_weights(...) | 39 |
| 23.2.2transfer_assisted_model_b_layer_weights(...) | 39 |
| 23.3Behavioral contract | 39 |
| 23.4Invariants | 40 |
| 24 Utility module: Evaluation | 40 |
| 24.1Purpose | 40 |
| 24.2Core utilities | 40 |
| 24.2.1evaluate_players(players, layout) | 40 |
| 24.2.2confidence_interval(mean, std_error) | 40 |
| 24.3Contract | 40 |
| 24.4Invariants | 40 |
| 25 Invariants | 40 |
| 25.1Purpose | 40 |
| 25.2Game and environment | 40 |
| 25.3Communication | 41 |
| 25.4Assisted / shared randomness | 41 |
| 25.5Trainable assisted models (Lin/Pyr) | 41 |
| 25.6Data generation and training | 41 |
| 25.7Logging and evaluation | 41 |
| 25.8Testability expectation (recommendation) | 42 |

1 QSeaBattle “ Design Specification

1.1 Purpose

This document is the **authoritative design specification** for the QSeaBattle project. It defines the game mechanics, player models, algorithms, training regimes, and utilities that together constitute the QSeaBattle system.

If any behavior in the codebase contradicts this specification, the code is considered incorrect.

1.2 Scope and audience

This specification is intended for:

- Developers implementing or modifying the QSeaBattle codebase
- Researchers reviewing the algorithmic and information-theoretic aspects
- Future maintainers re-implementing components independently

This is **not** a user tutorial; it is a technical design contract.

1.3 Structure of the specification

The document is organized into the following logical parts:

1. **Algorithms**
 - High-level, informative explanations of the classical assisted strategy, trainable Lin models, and hierarchical Pyr models.
2. **Conventions**
 - Global notation, terminology, and RFC-style contract language.
3. **Game infrastructure**
 - Game layout, environment, game execution, and tournaments.
4. **Players**
 - Abstract player interfaces and concrete player families (deterministic, assisted, neural).
5. **Shared randomness**
 - Classical shared-randomness resources and differentiable variants.
6. **Models**
 - Trainable assisted models (Lin and Pyr) with strict information constraints.
7. **Training**
 - Imitation learning, DRU/DIAL, and reinforcement learning regimes.
8. **Utilities**
 - Dataset generation, training helpers, evaluation, and weight transfer.
9. **Invariants**
 - One-sentence global truths that must always hold.

1.4 Normative vs informative content

- Chapters describing components, interfaces, and constraints are **normative**.
- The Algorithms chapter is **informative** and provides intuition only.
- In case of conflict, **normative chapters take precedence**.

Normative language uses: - MUST / MUST NOT - SHOULD / SHOULD NOT - MAY

1.5 Versioning and stability

This specification evolves with the codebase. Breaking changes to behavior MUST be accompanied by updates to this document.

A frozen PDF export may be tagged as: - *QSeaBattle Specification vX.Y*

1.6 How to read this document

- Start with **Algorithms** for intuition
- Read **Conventions** and **Invariants** carefully
- Then read the chapters relevant to the component you are working on

1.7 Repository layout (informative)

```
QSeaBattle/
├── docs/                                # This specification (Markdown source)
├── mkdocs.yml                            # Documentation build configuration
├── src/                                   # Implementation
├── tests/                                 # Tests (should reference invariants)
└── notebooks/                            # Experiments and analysis
```

1.8 Design principle (non-negotiable)

Information flow is the primary design constraint.
Learning may approximate logic, but must never bypass it.

2 Algorithms overview (informative)

2.1 Purpose

This chapter provides a consolidated, **informative (non-normative)** explanation of the core algorithms underlying QSeaBattle. It complements the normative chapters by explaining *how* the system works step by step, without introducing new requirements.

Nothing in this chapter overrides behavioral contracts defined elsewhere.

2.2 1. Classical assisted algorithm (baseline)

2.2.1 Problem

Player A observes a binary field of size n^2 .
Player B observes a one-hot gun position.
Only **one communication bit** is allowed.

2.2.2 High-level idea

Use **shared randomness** to correlate partial information held by the players so that Player B can reconstruct the relevant field bit with probability greater than any unassisted classical strategy.

2.2.3 Algorithm (conceptual)

1. Player A groups the field into pairs.
2. For each pair, Player A performs a shared-randomness measurement.
3. Outcomes reduce the field size by half.
4. Steps 1–3 repeat until a single bit remains.
5. Player A sends this bit to Player B.
6. Player B performs corresponding measurements in reverse order, guided by the gun position.

This forms a **pyramid reduction / reconstruction scheme**.

2.3 2. Shared randomness algorithm

2.3.1 Resource model

A shared-randomness resource produces correlated binary outcomes (x, y) such that:

$$[P(x = y) = p_{\text{high}}, \quad P(x \neq y) = 1 - p_{\text{high}}]$$

2.3.2 Usage rules

- Exactly one measurement per party
- First measurement is unbiased
- Second measurement is conditionally biased

This abstracts quantum-like correlations while remaining classical.

2.4 3. Deterministic baselines (Simple and Majority)

2.4.1 Why they exist

Deterministic baselines provide: - sanity checks (does the environment work?), - stable regression baselines, - and lower bounds for performance.

They do **not** use shared randomness.

2.4.2 Simple baseline (conceptual)

A minimal deterministic protocol: 1. Player A computes a fixed summary of the field (e.g., parity of a fixed subset) and sends it as comm. 2. Player B maps (gun, comm) through a fixed rule to decide shoot.

This baseline is intentionally simple and implementation-defined, but MUST remain deterministic.

2.4.3 Majority baseline (conceptual)

A stronger heuristic: 1. Player A sends whether the majority of a fixed subset S of the field is 1. 2. Player B uses this as its guess (optionally region-conditioned if comms_size > 1).

2.5 4. Lin trainable assisted algorithm

2.5.1 Goal

Approximate the classical assisted algorithm using minimal-capacity neural networks.

2.5.2 Algorithm

1. Replace fixed measurement logic with a linear measurement layer.
2. Replace fixed combination logic with a linear combine layer.
3. Train both layers using imitation learning from the classical policy.
4. Enforce all information-flow constraints at runtime.

The Lin model performs the same *logical* algorithm as the classical strategy, but learns the mapping from data.

2.6 5. Pyramid (Pyr) trainable assisted algorithm

2.6.1 Motivation

Lin models do not scale efficiently to large n^2 .

2.6.2 Algorithm

For $k = \log_2(n^2)$ layers:

Forward (Player A): 1. Apply a trainable measurement layer. 2. Perform one shared-randomness measurement. 3. Combine outputs to halve the field size.

Backward (Player B): 1. Use the gun position to select the relevant subtree. 2. Combine shared-randomness outcomes layer by layer. 3. Apply the communication bit at the final step.

The Pyr algorithm preserves logarithmic depth and strict information control.

2.7 6. Training algorithms (summary)

2.7.1 Imitation learning

1. Generate optimal trajectories using classical assisted players.
2. Train measurement and combine layers to match targets.
3. Validate that trained models obey runtime contracts.

2.7.2 Key restriction

Training may change **weights**, but never:
- Communication bandwidth
- Information access
- Layer structure

2.8 Relationship to the rest of the spec

| Chapter | Role |
|---------------------|------------------------|
| 3 | Game mechanics |
| 6 | Classical coordination |
| 7 | Trainable baseline |
| 8 | Hierarchical scaling |
| 9 | Safe utilities |
| This chapter | Algorithmic intuition |

2.9 Reader guidance

This chapter is intended for:
- Readers new to the project
- Reviewers evaluating novelty
- Future re-implementations

For authoritative behavior, always defer to the normative chapters.

3 Conventions

3.1 Purpose

This document defines **global notation**, **terminology**, and **cross-cutting conventions** used throughout the QSeaBattle specifica-

tion. Unless explicitly marked as *informative*, statements in this file are **normative**.

3.2 Notation

3.2.1 Sizes and shapes

- Let `field_size` = n (board side length). The total number of cells is n^2 .
- Vectors representing the full field or gun are **flattened** to length n^2 .
- Batch dimension is denoted by B .

Shape conventions - field: (B, n^2) or $(n^2,)$ for a single sample - gun: (B, n^2) or $(n^2,)$ one-hot - comm: (B, m) where $m = \text{comms_size}$ (often $m=1$) - shoot: $(B, 1)$ or scalar $\{0,1\}$

3.2.2 Binary values

Unless stated otherwise, binary arrays use values in $\{0,1\}$.

3.3 Terminology

3.3.1 Game vs tournament

- A **game** is one execution of the pipeline: generate field+gun → A decides → noise → B decides → reward.
- A **tournament** is a batch of games under identical layout parameters.

3.3.2 Players and models

- **Players** are decision-makers that implement `PlayerA.decide` and `PlayerB.decide`.
- **Models** are reusable computational components (e.g., neural models) that may back a player.
- A trainable model may store **per-decision state** (e.g., per-layer outcomes) required for Player B.

3.3.3 Shared randomness (SR)

- **Shared randomness** is a two-party resource used once per layer (Pyr) or once per decision (Lin).
- SR is **not communication**. SR does not increase `comms_size`.

3.4 Determinism, seeding, and reproducibility

- Any utility that samples randomness (data generation, evaluation) MUST accept a seed parameter or use a reproducible RNG strategy.
- Runtime gameplay MAY be stochastic due to SR and channel noise; evaluation runs SHOULD support deterministic replay via fixed seeds.

3.5 Channel noise

- Channel noise flips each bit of comm independently with probability channel_noise.
- Noise MUST be applied **after** Player A decides and **before** Player B decides.

3.6 Contract language (RFC-style)

The following keywords are normative: - **MUST / MUST NOT**: required for compliance - **SHOULD / SHOULD NOT**: recommended; deviations require justification - **MAY**: optional behavior permitted by the spec

3.7 Interface conventions

3.7.1 decide(...)

- decide methods MUST be pure with respect to their explicit inputs, except for documented internal state needed for coordination (e.g., per-layer outcomes).
- For comms_size == 1, comm MUST still be represented as an array/tensor of shape (B, 1) (not a scalar).

3.7.2 Dtypes

- Numpy arrays are acceptable for environment-level logic.
- TensorFlow tensors are acceptable for trainable models and training utilities.
- A given function MUST document whether it expects NumPy or Tensor inputs if ambiguous.

3.8 Error handling

- Shape mismatches MUST raise ValueError (preferred) rather than silently reshaping.

- Double-use of shared randomness MUST raise an error (typically `ValueError` or `RuntimeError`), not silently proceed.

3.9 Document status

- Chapters describing components are **normative**, unless explicitly labeled *informative*.
- `docs/algorithms.md` is **informative** and does not override contracts.

4 Class GameLayout

4.1 Purpose

Define the immutable parameters that specify a QSeaBattle game configuration, ensuring consistency across environment, players, tournaments, and logs.

4.2 Location

- **Module:** `src/Q_Sea_Battle/game_layout.py`
- **Class:** `GameLayout`

4.3 Public attributes

| Name | Type | Constraints | Description |
|--------------------------------------------|------------------------|---------------|----------------------------------------------|
| <code>field_size</code> | <code>int</code> | power of 2 | Board dimension n; total cells n^2 |
| <code>comms_size</code> | <code>int</code> | divides n^2 | Number of communication bits |
| <code>enemy_probability</code> | <code>float</code> | [0,1] | Probability a field cell equals 1 |
| <code>channel_noise</code> | <code>float</code> | [0,1] | Bit-flip probability on comm channel |
| <code>number_of_games_in_tournament</code> | <code>int</code> | >0 | Games per tournament |
| <code>log_columns</code> | <code>list[str]</code> | — | Columns stored in <code>TournamentLog</code> |

4.4 Public methods

4.4.1 `from_dict(parameters: dict) -> GameLayout`

Create a validated GameLayout from a dictionary.

4.4.2 `to_dict() -> dict`

Return all layout parameters as a dictionary.

4.5 Invariants

- `field_size` MUST be a power of two.
- `comms_size` MUST divide `field_size^2`.
- All components MUST share the same GameLayout instance.

4.6 Failure modes

- `ValueError` if constraints are violated.

5 Class GameEnv

5.1 Purpose

Implement the physical rules of QSeaBattle: field generation, gun placement, reward evaluation, and communication noise.

5.2 Location

- **Module:** `src/Q_Sea_Battle/game_env.py`
- **Class:** `GameEnv`

5.3 Public methods

5.3.1 `reset() -> None`

Generate a new random field and gun.

5.3.2 `provide() -> tuple[np.ndarray, np.ndarray]`

Return (field, gun) as flattened binary arrays.

5.3.3 `evaluate(shoot: int) -> float`

Return 1.0 if decision matches field value at gun index, else 0.0.

5.3.4 `apply_channel_noise(comm: np.ndarray) -> np.ndarray`

Apply independent bit flips with probability `channel_noise`.

5.4 Invariants

- `gun` MUST be one-hot.
- `field.shape == gun.shape == (n^2,)`

6 Class Game

6.1 Purpose

Execute a single QSeaBattle game.

6.2 Method

6.2.1 `play()`

Runs one game: reset, decide A, apply noise, decide B, evaluate.

7 Class Tournament

7.1 Purpose

Run multiple games and collect structured data.

7.1.1 `tournament() -> TournamentLog`

Run a full tournament.

8 Class TournamentLog

8.1 Purpose

Store per-game data for analysis and training.

8.1.1 `Methods`

- `update(...)`
- `outcome() -> (mean, std_error)`

9 Base player interfaces

9.1 Purpose

Define the abstract interfaces for Player A and Player B, plus the factory interface that returns a matched pair. Unless otherwise specified by a concrete implementation, the **default baseline behavior is random**.

9.2 Class Players

9.2.1 Purpose

Factory and lifecycle manager for paired PlayerA / PlayerB instances.

9.2.2 Public methods

- `players() -> tuple[PlayerA, PlayerB]`
Return a matched (player_a, player_b) pair.
- `reset() -> None`
Reset any internal per-game state (if applicable).

9.2.3 Invariants

- `players()` MUST always return a matched A/B pair.
- Both players MUST share the same GameLayout.

9.3 Class PlayerA

9.3.1 Purpose

Observe the field and produce a communication vector.

9.3.2 Method

`decide(field: np.ndarray, supp: Any | None = None) -> np.ndarray`

9.3.3 Contract

- Output MUST be a binary array of length `comms_size`.
- Even when `comms_size == 1`, output MUST be an array (shape `(1,)` or `(B,1)`), not a scalar.
- Player A MUST NOT access gun information.

9.3.4 Default baseline behavior (random)

If a concrete implementation does not specify a different decision rule, PlayerA.decide is interpreted as: - Sample each communication bit independently and uniformly from {0,1}.

This provides a well-defined **random baseline** for debugging and sanity checks.

9.4 Class PlayerB

9.4.1 Purpose

Observe the gun position and received communication, and decide whether to shoot.

9.4.2 Method

```
decide(gun: np.ndarray, comm: np.ndarray, supp: Any | None = None) -> int
```

9.4.3 Contract

- Output MUST be scalar 0 or 1 (or (B,1) for batched implementations).
- Input gun MUST be one-hot.
- Player B MUST NOT access the full field.

9.4.4 Default baseline behavior (random)

If a concrete implementation does not specify a different decision rule, PlayerB.decide is interpreted as: - Sample shoot uniformly from {0,1}, independent of (gun, comm).

9.5 Related

- Implemented by: deterministic, assisted, neural players
- Algorithms: docs/algorithms.md (informative)

10 Deterministic and baseline players

10.1 Purpose

Define simple, non-learning baseline player families used for: - sanity checks, - debugging, - performance lower bounds, - and validating tournament logging.

This chapter specifies two baseline families: 1. **SimplePlayers** (a minimal baseline; Player B includes a stochastic fallback) 2. **MajorityPlayers** (a deterministic heuristic baseline)

10.2 Class SimplePlayers(Players)

10.2.1 Purpose

Factory that returns a matched (SimplePlayerA, SimplePlayerB) pair.

10.2.2 Behavioral contract

- SimplePlayers.players() MUST generate instances of SimplePlayerA and SimplePlayerB.
- The family MUST respect communication bandwidth: output length equals layout.comms_size.

10.3 Class SimplePlayerA(PlayerA)

10.3.1 Purpose

Send the values of the first m field cells directly, where $m = \text{layout.comms_size}$.

10.3.2 Algorithm (normative)

Let $m = \text{comms_size}$. Let field be the flattened field vector of length n^2 .

1. Read the first m entries: $\text{field}[0:m]$.
2. Output $\text{comm} = \text{field}[0:m]$.

10.3.3 Invariants

- MUST NOT access gun information.
- MUST output exactly m bits.

10.4 Class SimplePlayerB(PlayerB)

10.4.1 Purpose

If the gun points at one of the first m cells, use the communicated value for that cell. Otherwise, fall back to a simple prior-based stochastic guess using $\text{layout.enemy_probability}$.

10.4.2 Algorithm (normative)

Let $m = \text{comms_size}$. Let gun be one-hot of length n^2 . Let $\text{idx} = \text{argmax}(\text{gun})$.

1. If $\text{idx} < m$:
 - Output $\text{shoot} = \text{comm}[\text{idx}]$.
2. Else:
 - Sample $\text{shoot} \sim \text{Bernoulli}(\text{layout.enemy_probability})$.

10.4.3 Notes

This baseline is **not fully deterministic** due to the stochastic fallback in step 2.

10.4.4 Invariants

- MUST NOT access full field information.
- Output MUST be binary (0 or 1).
- If $\text{idx} < m$, the output MUST equal the communicated bit for that index.

10.5 Class MajorityPlayers(Players)

10.5.1 Purpose

Factory that returns a matched (`MajorityPlayerA`, `MajorityPlayerB`) pair implementing a majority heuristic over partitions.

10.5.2 Behavioral contract

- `MajorityPlayers.players()` MUST generate instances of `MajorityPlayerA` and `MajorityPlayerB`.
- The family MUST respect communication bandwidth: output length equals `layout.comms_size`.
- The family MUST be deterministic given inputs.

10.6 Class MajorityPlayerA(PlayerA)

10.6.1 Purpose

Split the flattened field into m pieces and send one majority bit per piece.

10.6.2 Algorithm (normative)

Let $m = \text{comms_size}$. Let $N = n^2$ be the flattened field length.

1. Partition indices $\{0, 1, \dots, N-1\}$ into m contiguous blocks:
 - Block j contains indices from $\text{start} = \text{floor}(j*N/m)$ up to (but excluding) $\text{end} = \text{floor}((j+1)*N/m)$.
2. For each block j :
 - Let $\text{ones}_j = \sum(\text{field}[i] \text{ for } i \text{ in block } j)$.
 - Let $\text{zeros}_j = |\text{block } j| - \text{ones}_j$.
 - Set $\text{comm}[j] = 1$ if $\text{ones}_j \geq \text{zeros}_j$, else $\text{comm}[j] = 0$.
3. Output comm (length m).

10.6.3 Invariants

- MUST NOT access gun information.
- MUST output exactly m bits.
- MUST be deterministic given field.

10.7 Class MajorityPlayerB(PlayerB)

10.7.1 Purpose

Determine which subset (block) the gun points to and return the corresponding communicated majority bit.

10.7.2 Algorithm (normative)

Let $m = \text{comms_size}$. Let $N = n^2$. Let $\text{idx} = \text{argmax}(\text{gun})$.

1. Compute the block id j such that idx lies in block j under the same partition rule used by MajorityPlayerA.
2. Output $\text{shoot} = \text{comm}[j]$.

10.7.3 Invariants

- MUST NOT access full field information.
- Output MUST be binary.
- MUST be deterministic given (gun , comm).

10.8 Related

- Base interfaces: [docs/players/base.md](#)
- Game parameters: [docs/game/GameLayout.md](#)

11 Assisted players (shared-randomness assisted)

11.1 Purpose

Define Player A and Player B strategies that use a shared-randomness resource to outperform unassisted baselines under strict communication constraints.

This chapter specifies the assisted-player family and its dependency on the shared randomness contract.

11.2 Shared randomness dependency

Assisted players MUST use a resource that satisfies:

- docs/shared_randomness/shared_randomness.md (normative)

The resource may be implemented as: - a pure Python class (`SharedRandomness`), or - a differentiable Keras layer (`SharedRandomnessLayer`), but the externally observable behavior MUST comply with the same contract.

11.3 Class SharedRandomness()

11.3.1 Purpose

Reference class name for the shared randomness resource used by assisted players.

The authoritative behavior is defined in `docs/shared_randomness/shared_randomness.md`

11.4 Class AssistedPlayers(Players)

11.4.1 Purpose

Factory and manager for assisted (`AssistedPlayerA`, `AssistedPlayerB`) pairs that share correlated resources.

11.4.2 Location

- **Module:** `src/Q_Sea_Battle/assisted_players.py`
- **Class:** `AssistedPlayers`

11.4.3 Parameters

| Name | Type | Constraints |
|--------|------------|----------------------------------------|
| p_high | float | [0,1] |
| layout | GameLayout | comms_size == 1 (recommended baseline) |

11.4.4 Behavioral contract

- Communication bandwidth MUST be respected (comms_size bits).
- Assisted players MUST NOT treat shared randomness as extra communication.
- Per game, the shared randomness resources MUST be reset.

11.5 Class AssistedPlayerA(PlayerA)

11.5.1 Purpose

Observe the field and compute comm using shared randomness according to a fixed hierarchical reduction scheme.

11.5.2 Behavioral contract (normative)

- MUST NOT access gun information.
- MUST consume shared randomness in the same sequence that Player B will mirror.
- MUST produce exactly comms_size bits (typically 1).

11.5.3 Algorithmic outline (informative)

- Reduce the field iteratively (pairwise grouping) while consuming one SR measurement per reduction stage.
- Store per-stage outcomes needed for Player B reconstruction.

11.6 Class AssistedPlayerB(PlayerB)

11.6.1 Purpose

Use gun position, received comm, and shared randomness outcomes to reconstruct the relevant field bit.

11.6.2 Behavioral contract (normative)

- MUST NOT access full field information.
- MUST mirror Player A's SR consumption order (one SR measurement per stage).

- Gun input MUST remain one-hot throughout any internal transformations.
- Output MUST be binary at inference time.

11.6.3 Algorithmic outline (informative)

- Traverse the reduction hierarchy guided by gun position.
- Combine SR outcomes and comm to compute shoot.

11.7 Invariants

- For the standard assisted protocol: `comms_size == 1` and n^2 is a power of two.
- Each SR resource is measured at most once per party per game.

12 Neural players (unassisted)

12.1 Purpose

Define unconstrained neural-network-based players that do not use shared randomness. These players serve as comparison baselines against assisted strategies.

12.2 Characteristics

- Communication limited to `comms_size`
- No shared randomness
- Fully trainable via reinforcement learning or supervised learning

12.3 Behavioral constraints

- Player A MUST NOT access gun information
- Player B MUST NOT access full field information
- Communication bandwidth MUST be respected

12.4 Notes

Neural players may violate classical assisted performance bounds, but serve as useful empirical baselines.

13 Shared randomness resource (non-classical)

13.1 Purpose

Define a two-party shared-randomness resource that produces **non-classical correlations** between Player A and Player B outcomes, without increasing communication bandwidth.

This specification defines the **behavioral contract** of shared randomness. Multiple implementations (e.g., a pure Python class and a Keras layer) MUST satisfy the same contract.

13.2 Location

- **Python implementation (example):** `src/Q_Sea_Battle/shared_randomness.py`
 ↳ class `SharedRandomness`
- **Differentiable implementation (example):** `src/Q_Sea_Battle/shared_randomness_layer.py`
 ↳ class `SharedRandomnessLayer`

13.3 Parameters

| Name | Type | Constraints | Description |
|---------------------|--------------------|-------------|------------------------------------------------------------|
| <code>length</code> | <code>int</code> | > 0 | Number of correlated outcome bits produced per measurement |
| <code>p_high</code> | <code>float</code> | $[0, 1]$ | Correlation strength parameter |

13.4 Public interface (conceptual)

13.4.1 `reset() -> None`

Reset per-game state so the resource can be used again.

13.4.2 `measurement_a(measurement: array_like) -> np.ndarray`

Player A performs a measurement and receives a binary outcome vector of length `length`.

13.4.3 `measurement_b(measurement: array_like) -> np.ndarray`

Player B performs a measurement and receives a binary outcome vector of length `length`.

The measurement input is a *measurement choice* (a setting) and MAY be encoded as a binary vector, integer index, or tensor but it MUST be documented by the implementation and used consistently by both parties.

13.5 Behavioral contract (normative)

13.5.1 Single-use per party

- Each party MUST call its measurement method at most once per resource instance per game.
- Calling `measurement_a` twice, or `measurement_b` twice, MUST raise an error.

13.5.2 Two-measurement lifecycle

- The resource supports up to two measurements total (one per party) between resets.

13.5.3 Output format

- Outcomes MUST be binary arrays in $\{0,1\}$ of shape $(\text{length},)$ (or (B, length) for batched/tensor implementations).

13.5.4 Non-classical correlation requirement

Let x be Player A's outcome, y Player B's outcome (bitwise).

The joint distribution MUST satisfy:

- $P(x_i = y_i) = p_{\text{high}}$
- $P(x_i \neq y_i) = 1 - p_{\text{high}}$

for each bit index i , **conditional on both parties using the same measurement choice**.

If measurement choices differ, the implementation MUST still: - return valid binary outcomes, and - produce correlations that are a documented function of both measurement choices.

This clause is what makes the resource non-classical in this project: the correlation structure is an explicit, tunable contract, and not restricted to classical shared coins.

13.5.5 No information leakage

- The resource MUST NOT depend on field, gun, or any game secrets.
- It MUST NOT encode additional side channels beyond its specified outputs.

13.6 Failure modes

- ValueError (recommended) if a party measures twice.
- RuntimeError if used without reset across games (implementation-dependent).

13.7 Notes (informative)

Two compliant implementations are common: 1. **Pure Python**: explicit sampling of correlated bits with the required joint distribution. 2. **Keras layer**: a differentiable proxy used during training that matches the same marginal/correlation behavior.

14 SharedRandomnessLayer (differentiable implementation)

14.1 Purpose

Provide a differentiable implementation of the **Shared randomness resource (non-classical)** contract, to enable gradient-based training (imitation, DRU/DIAL, RL variants) while preserving runtime semantics.

This layer is an **implementation** of the shared randomness contract; it does not define new behavior.

14.2 Location

- **Module:** src/Q_Sea_Battle/shared_randomness_layer.py
- **Class:** SharedRandomnessLayer

14.3 Inputs / Outputs

Implementations MAY support:
- Unbatched (length,) inputs and outputs
- Batched (B, length) inputs and outputs
- TensorFlow tensors throughout

Outcomes MUST be interpretable as binary at runtime:
- During training: may use continuous relaxations (e.g., logits, straight-through estimators)
- During inference/evaluation: MUST produce discrete {0,1} outcomes (or be replaced by Python SharedRandomness)

14.4 Behavioral contract

The layer MUST satisfy all requirements from: - docs/shared_randomness/shared_randomness.

Additionally:
- The differentiable relaxation MUST NOT increase the number of communicated bits.
- Any temperature/annealing schedule MUST be documented and controlled by training utilities.

14.5 Recommended modes (informative)

- mode="train": differentiable relaxation (logits / sigmoid / straight-through)
- mode="eval": discrete sampling consistent with p_high

14.6 Failure modes

- Shape mismatch MUST raise ValueError.

15 Trainable assisted models â€” Linear (Lin)

15.1 Purpose

Provide a trainable approximation of the classical assisted strategy using linear neural network layers, preserving all behavioral contracts defined for assisted players.

15.2 Scope

This chapter specifies the **Lin** family of trainable assisted models. These models replace fixed logical operations with learned linear mappings while maintaining the same information-theoretic constraints.

15.3 Class LinTrainableAssistedModelA

15.4 Purpose

Approximate the classical AssistedPlayerA strategy using a single linear measurement layer and a linear combine layer.

15.5 Location

- **Module:** src/Q_Sea_Battle/lin_trainable_assisted_models.py
- **Class:** LinTrainableAssistedModelA

15.6 Inputs

| Name | Shape | Description |
|-------|----------------------|---------------------|
| field | (B, n ²) | Binary field vector |

15.7 Outputs

| Name | Shape | Description |
|------|--------|-------------------|
| comm | (B, 1) | Communication bit |

15.8 Internal state

- `measurements_per_layer`: list of tensors
- `outcomes_per_layer`: list of tensors

15.9 Behavioral contract

- Exactly **one measurement** is performed per decision.
- Exactly **one communication bit** is produced.
- The model **MUST NOT** use gun information.
- Shared randomness is accessed via a differentiable proxy.

15.10 Invariants

- Input length n^2 **MUST** be a power of two.
- `comms_size == 1`.

15.11 Class LinTrainableAssistedModelB

15.12 Purpose

Approximate the classical AssistedPlayerB strategy using linear layers to reconstruct the relevant field bit.

15.13 Location

- **Module:** src/Q_Sea_Battle/lin_trainable_assisted_models.py
- **Class:** LinTrainableAssistedModelB

15.14 Inputs

| Name | Shape | Description |
|------|----------------------|--------------------|
| gun | (B, n ²) | One-hot gun vector |
| comm | (B, 1) | Communication bit |

15.15 Outputs

| Name | Shape | Description |
|-------|--------|-----------------|
| shoot | (B, 1) | Binary decision |

15.16 Behavioral contract

- Exactly **one measurement** is performed per decision.
- Output **MUST** be binary at inference time.
- The model **MUST NOT** access the field directly.

15.17 Invariants

- Gun input **MUST** be one-hot.
- Decision depends only on (gun, comm, shared randomness).

15.18 Notes / rationale

The Lin models serve as the minimal trainable baseline. They are intentionally capacity-limited to: - Test learnability under strict information constraints - Provide a reference point for deeper architectures (Pyr)

All violations of classical-assisted invariants are considered specification errors.

16 Trainable assisted models â€” Pyramid (Pyr)

16.1 Purpose

Define a hierarchical, multi-layer trainable assisted model that scales the classical assisted strategy to large fields by iteratively reducing the problem size while preserving all assisted-player contracts.

The Pyramid (Pyr) models generalize the Lin models by composing repeated measurementâ€“combine stages in a strictly structured hierarchy.

16.2 Conceptual overview

Let the field have size n^2 , with $n^2 = 2^k$. The Pyr architecture performs k iterations.

At each iteration: - The effective field size is halved - One shared-randomness measurement is consumed - Intermediate outcomes are stored for later reconstruction

Exactly one communication bit is produced at the final layer.

16.3 Class PyrTrainableAssistedModelA

16.4 Purpose

Approximate the classical AssistedPlayerA pyramid strategy using trainable neural layers, producing a single communication bit from a large field.

16.5 Location

- **Module:** src/Q_Sea_Battle/pyr_trainable_assisted_models.py
- **Class:** PyrTrainableAssistedModelA

16.6 Inputs

| Name | Shape | Description |
|-------|----------------------|---------------------|
| field | (B, n ²) | Binary field vector |

16.7 Outputs

| Name | Shape | Description |
|------|--------|-------------------|
| comm | (B, 1) | Communication bit |

16.8 Internal state

The model **MUST** store the following per decision:

- `measurements_per_layer[i]`: measurement tensor at layer i
- `outcomes_per_layer[i]`: shared-randomness outcomes at layer i

where: - Layer 0 operates on size n^2 - Layer $i + 1$ operates on size $n^2/2^{i+1}$

16.9 Iterative structure (normative)

For layer $i = 0, \dots, k - 1$:

1. **Measurement layer**
 - Input size: $n^2/2^i$
 - Output size: $n^2/2^{i+1}$
2. **Shared randomness**
 - One measurement is performed
 - Outcome size: $n^2/2^{i+1}$
3. **Combine layer**
 - Inputs: measurement output + shared randomness
 - Output: reduced field for next layer

At the final layer, the reduced field has size 1 and is emitted as `comm`.

16.10 Behavioral contract

- Exactly **one shared-randomness measurement per layer**
- Exactly **one communication bit total**
- No access to gun information
- No skipping or reordering of layers

16.11 Invariants

- Input size **MUST** be a power of two
- Number of layers = $\log_2(n^2)$
- Lengths of `measurements_per_layer` and `outcomes_per_layer` **MUST** equal number of layers

16.12 Failure modes

- `ValueError` if input size is not a power of two
- `RuntimeError` if layers are executed inconsistently

16.13 Class PyrTrainableAssistedModelB

16.14 Purpose

Reconstruct the relevant field bit using hierarchical outcomes, gun position, and a single communication bit.

16.15 Location

- **Module:** `src/Q_Sea_Battle/pyr_trainable_assisted_models.py`
- **Class:** `PyrTrainableAssistedModelB`

16.16 Inputs

| Name | Shape | Description |
|------|----------------------|--------------------|
| gun | (B, n ²) | One-hot gun vector |
| comm | (B, 1) | Communication bit |

16.17 Outputs

| Name | Shape | Description |
|-------|--------|-----------------|
| shoot | (B, 1) | Binary decision |

16.18 Internal state

- `outcomes_per_layer[i]`: outcomes received from Player A
- Optional learned reconstruction layers per hierarchy level

16.19 Reconstruction procedure (normative)

For layer $i = k - 1, \dots, 0$:

1. Use gun position to select the relevant subtree
2. Combine:
 - Corresponding shared-randomness outcome
 - Current reconstruction state
 - Communication bit (only at final stage)

The final output is a scalar decision.

16.20 Behavioral contract

- Exactly **one shared-randomness measurement per layer**
- Gun input **MUST** remain one-hot throughout reconstruction
- Output **MUST** be binary at inference time

16.21 Invariants

- Reconstruction depth equals measurement depth
- No access to full field information

16.22 Notes / rationale

The Pyr architecture:
- Scales logarithmically with field size
- Mirrors the classical pyramid protocol
- Makes information flow explicit and inspectable

Any deviation from the layer structure constitutes a **spec violation**.

17 Imitation learning for assisted models

17.1 Purpose

Train Lin (and later Pyr) assisted models to imitate the optimal classical assisted strategy using supervised learning.

17.2 Data generation

- Generate (field, gun) samples using GameEnv
- Generate optimal (comm, shoot) targets using classical AssistedPlayers

17.3 Behavioral contract

- Training data **MUST** be generated from a spec-compliant classical policy.
- Training **MUST NOT** introduce additional information channels.

17.4 Losses

- Binary cross-entropy for communication bit
- Binary cross-entropy for shooting decision

17.5 Invariants

- Dataset inputs and targets **MUST** match the shapes defined in Chapter 7.
- Training **MUST NOT** alter the runtime contracts of the models.

18 DRU / DIAL training

18.1 Purpose

Specify training using Discretize–Regularize Units (DRU) and Differentiable Inter-Agent Learning (DIAL) for communication-constrained agents.

18.2 Algorithm (informative)

- Replace hard binary communication with differentiable proxies
- Regularize outputs toward discrete values
- Anneal regularization during training

18.3 Constraints

- Runtime communication **MUST** remain discrete
- DRU/DIAL **MUST NOT** increase communication bandwidth
- Learned policies **MUST** satisfy all player contracts

18.4 Notes

DRU/DIAL is optional and experimental in QSeaBattle.

19 Reinforcement learning

19.1 Purpose

Specify reinforcement-learning-based training regimes for players, primarily for baseline comparison.

19.2 Supported approaches

- Policy gradients
- Self-play
- Actor–critic methods

19.3 Constraints

- RL MUST respect all information-flow constraints
- Reward signals MUST come exclusively from GameEnv
- Exploration MUST NOT leak additional information

19.4 Notes

Reinforcement learning is not the primary focus of QSeaBattle, but provides useful empirical benchmarks.

20 Training considerations for Pyramid (Pyr) models

20.1 Purpose

Specify additional constraints and recommendations for training Pyramid assisted models.

20.2 Training regimes

- Supervised imitation (from classical pyramid policy)
- Layer-wise pretraining (optional)
- End-to-end fine-tuning (allowed but constrained)

20.3 Behavioral constraints

- Training MUST NOT introduce cross-layer shortcuts
- Shared-randomness usage MUST remain one-per-layer
- Communication bandwidth MUST remain 1 bit

20.4 Diagnostics (recommended)

- Per-layer loss monitoring
- Consistency checks between adjacent layers
- Ablation of individual layers

20.5 Invariants

- Runtime behavior MUST match Chapter 8 contracts
- Learned weights MUST NOT change input/output interfaces

21 Utility module: Data generation

21.1 Purpose

Provide deterministic, spec-compliant utilities for generating datasets used in training and evaluation, without introducing any additional information channels.

21.2 Scope

These utilities are used exclusively for **offline data generation**. They MUST NOT be used at runtime during actual games or tournaments.

21.3 Core utilities

21.3.1 generate_measurement_dataset_a(layout, num_samples, seed)

Generate training samples for Player A measurement layers.

Inputs - layout: GameLayout - num_samples: int - seed: int

Outputs - Dataset of (field, meas_target)

Contract - Targets MUST be generated using classical Assisted-PlayerA - No gun information is accessible

21.3.2 generate_measurement_dataset_b(layout, num_samples, seed)

Generate training samples for Player B measurement layers.

Contract - Targets MUST be generated using classical Assisted-PlayerB - Field information MUST NOT be accessible

21.4 Invariants

- Dataset shapes MUST match model interfaces
- Random seeds MUST make generation reproducible

22 Utility module: Imitation training

22.1 Purpose

Define helper utilities to train trainable assisted models using supervised imitation learning, while preserving all runtime behavioral contracts.

22.2 Core utilities

22.2.1 `train_layer(layer, dataset, loss, epochs, metrics=None)`

Train a single neural layer in isolation.

Contract - The utility MUST NOT alter layer input/output signatures
- The utility MUST NOT introduce cross-layer information flow

22.2.2 `train_model(model, datasets)`

Orchestrate training of full Lin or Pyr models.

Contract - Training MUST respect per-layer constraints (Chapter 7, 8)
- Runtime behavior MUST remain unchanged after training

22.3 Invariants

- Training utilities MUST NOT modify game logic
- All randomness MUST be controlled and logged

23 Utility module: Layer and weight transfer

23.1 Purpose

Provide safe utilities for transferring trained weights from standalone layers into composite assisted models.

23.2 Core utilities

23.2.1 `transfer_assisted_model_a_layer_weights(...)`

Copy trained Player A layers into a composite model.

23.2.2 `transfer_assisted_model_b_layer_weights(...)`

Copy trained Player B layers into a composite model.

23.3 Behavioral contract

- Only weights may be transferred
- Model topology MUST remain unchanged
- Any shape mismatch MUST raise an error

23.4 Invariants

- Post-transfer model MUST satisfy original spec contracts

24 Utility module: Evaluation

24.1 Purpose

Provide helpers for evaluating trained players using tournaments and aggregating results.

24.2 Core utilities

24.2.1 `evaluate_players(players, layout)`

Run a tournament and return performance statistics.

24.2.2 `confidence_interval(mean, std_error)`

Compute confidence intervals for reported results.

24.3 Contract

- Evaluation MUST NOT alter player state
- Evaluation MUST be reproducible given fixed seeds

24.4 Invariants

- Results MUST correspond to exactly the games played

25 Invariants

25.1 Purpose

This page lists **global invariants**: one-sentence truths that MUST hold across the QSeaBattle codebase. These are intended to be: - easy to audit during code review, - easy to test, - and stable over time.

25.2 Game and environment

- The field and gun vectors MUST have identical shape (n^2 ,) (or (B, n^2) batched).
- The gun vector MUST be one-hot: `gun.sum() == 1`.

- Channel noise MUST be applied only to comm, never to field or gun.
- Rewards MUST be deterministic given (field, gun, shoot).

25.3 Communication

- Communication bandwidth is exactly comms_size bits per game.
- If comms_size == 1, communication MUST still be represented as shape (B, 1) (never a scalar).

25.4 Assisted / shared randomness

- Shared randomness MUST NOT be treated as an extra communication channel.
- Each shared-randomness resource MAY be measured at most once by each party.
- For pyramid strategies, exactly one shared-randomness measurement is consumed per layer per party.

25.5 Trainable assisted models (Lin/Pyr)

- Player A models MUST NOT access gun information.
- Player B models MUST NOT access the full field information.
- Any learned model MUST preserve the same input/output interfaces as the classical policy it approximates.
- Model state required for reconstruction (e.g., per-layer outcomes) MUST be captured during Player A's decision and used by Player B; it MUST NOT be recomputed using hidden information.

25.6 Data generation and training

- Imitation-learning targets MUST be produced by a spec-compliant classical policy.
- Training utilities MUST NOT introduce new information paths (no shortcuts, no extra inputs).
- Weight transfer utilities MUST NOT change topology; only parameter values may be copied.

25.7 Logging and evaluation

- Each tournament log row MUST correspond to exactly one game.

- Evaluation utilities MUST NOT mutate player behavior beyond documented resets.

25.8 Testability expectation (recommendation)

- Every invariant above SHOULD have at least one unit test or property test that would fail if the invariant is violated.